



GBDi

Government Big Data Institute

สถาบันส่งเสริมการวิเคราะห์และบริหารข้อมูลขนาดใหญ่ภาครัฐ (สวช.)




Introduction to Supervised Learning: Basic Classification Models

Patipan Prasertsom

Classification

- In machine learning, classification problem is a problem of predicting the correct label (class) that a given data belongs to.
- There are many types of classification algorithms, such as
 - Regression-based classifiers: ex. logistic regressions
 - Tree-based classifiers: ex. decision tree, random forest, gradient tree boosting
 - Probability-based classifiers: ex. Bayes classifier
 - Similarity-based classifier: ex. k-Nearest Neighbor classifier
 - etc.
- Today, we will touch on some of the well-known classifiers



k-Nearest Neighbors classifier (kNN)

- **The idea:** similar data points are likely to be of the same type.
- We infer a class of a data point from its **k most similar** data
- How do we find the “most similar” (nearest) data points?
 - There are many way of measuring the distance between data point. One frequently discussed method is the Euclidean distance.
 - Euclidean distance – a distance between 2 points in cartesian coordinates

$$\text{Euclidean Distance} = \sqrt{(x_i - x_j)^2 + (y_i - y_j)^2}$$

- Basically asking, “**By looking at k data points that are most similar to me, which class am I most likely to be in?**”




Figure from Illuyanka
(https://commons.wikimedia.org/wiki/File:Dot_Product.svg)

k-Nearest Neighbors classifier (kNN)





Figure 6-2. Nearest neighbor classification. The point to be classified, labeled with a question mark, would be classified + because the majority of its nearest (three) neighbors are +.



- Observes k closest data point and decide the class of data points by voting.
- Usually, k is chosen as an odd number

- The choice of k matters!
- Different number of k can result in different predictions
- Feature scale matters!

k-Nearest Neighbors classifier (kNN)




Figure 6-4. Classification boundaries created on a three-class problem created by 1-NN (single nearest neighbor).

kNN models with a small k

- A finer granularity separation boundaries
- More susceptible to the presents of outlier.




Figure 6-5. Classification boundaries created on a three-class problem created by 30-NN (averaging 30 nearest neighbors).

kNN models with a large k


- More tolerant to noise
- Smooth but coarser separation boundaries

kNN Example

- Colab!

Logistic Regression

- Logistic regression is a binary (0/1) classification model, meaning it predict whether a data is of a certain class or not.
- **Motivation:** Let's say our data has classes ($y=0$ and $y=1$) as shown on the right. Can we use linear regression to predict them?
- While this looks like it can *somewhat* do the job, it is not appropriate
- We do not want the predicted output to be infinitely high/low. We simply want it to tell us whether the data is of our target type .
- Even better if it tells us how *confident* it is that the data is of our target type.





Logistic Regression

- A *sigmoid* function has a value ranging from 0 to 1 for its entire input range

$$y = \frac{1}{1 + e^{-x}}$$

- This makes it's a more appropriate choice to use compared to using purely linear equation



Logistic Regression


- Specifically, a logistic regression model uses a sigmoid function in conjunction with the linear equation.

$$y = \frac{1}{1 + e^{-f(x)}}$$

where


$$f(x) = w_0 + w_1x_1 + w_2x_2 + \dots + w_mx_m$$

- Fitting a logistic equation is akin to fitting a linear equation model which will then predict where we are on the input of the sigmoid function.
- This is similar to predicting how confidence we are of our prediction



Logistic Loss Function

- Since the goal for a logistic regression is to classify the data, its loss function needs to minimize the misclassification
 - Needs to have high value when predict incorrectly and low when predict correctly
- For example, we want a function that behaves like
 - $-\log(\hat{y}^{(i)})$ when $y^{(i)} = 1$
 - $-\log(1 - \hat{y}^{(i)})$ when $y^{(i)} = 0$



Logistic Loss Function

- The logistic loss function is

$$\text{Logistic Loss} = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log(\hat{y}^{(i)}) + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}))$$

where


$$\hat{y}^{(i)} = \frac{1}{1 + e^{-f(x^{(i)})}}$$

and

$$f(x^{(i)}) = w_0 + w_1 x_1^{(i)} + w_2 x_2^{(i)} + \dots + w_m x_m^{(i)}$$

- Optimization technique like [gradient descent](#) can then be applied to find the optimal set of parameters
- Also note that Logistic regression has linear classification boundary

Logistic Regression: Separation Boundary and Threshold



Logistic Regression

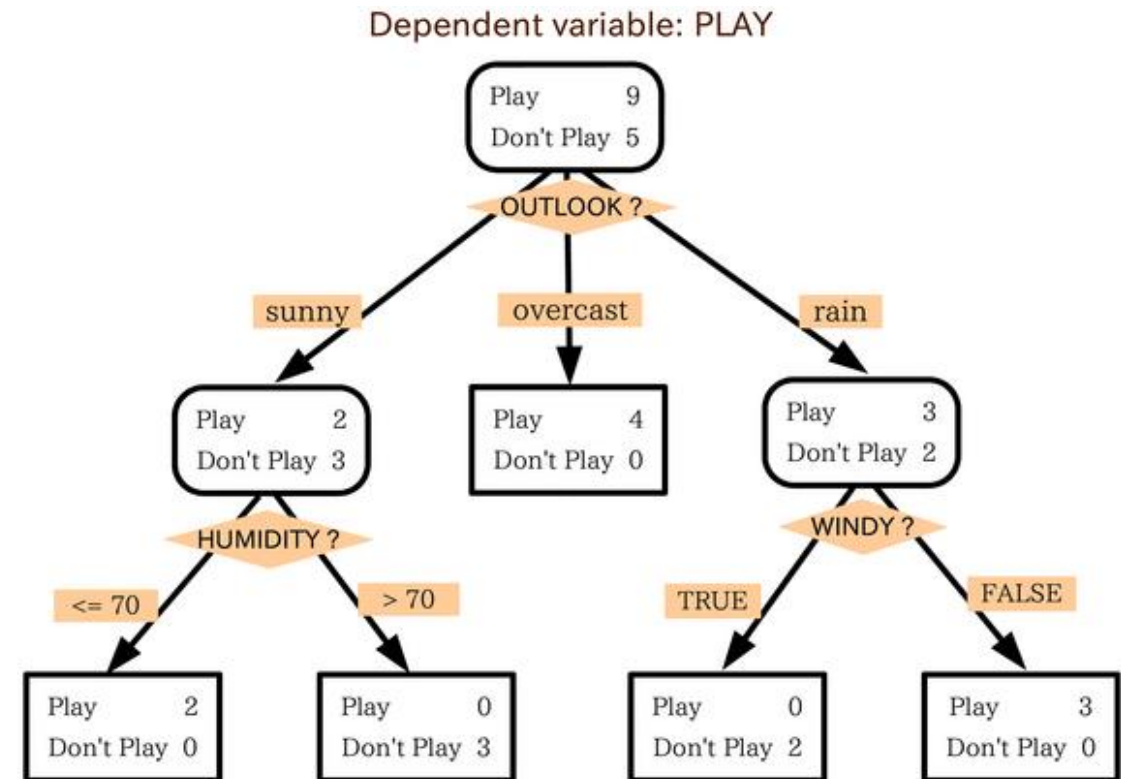
- One of the biggest advantage of logistic regression model is its **explainability**
- Similar to linear regression, the coefficients obtained for each feature in the final equation of the logistic regression are indicators of the **direction (positive/negative) and strength of the relationships of each feature and the prediction target**
 - Furthermore, this allows us to explain the expected consequence when changing the value of a feature.
- The coefficients obtained can also be used to perform feature selection by weeding out features with small coefficients.
- Also just like linear regression, regularization techniques can also be applied to improve the model's performance (in fact, regularization is applied by default in sklearn library)

Logistic Regression Example

- Colab!

Decision Tree

- A decision tree is a tree-like structure where the internal nodes represents series of “decisions” which eventually leads to an outcome (represented by the terminal nodes)
- Each node within a tree (apart from the final level) is a *segmentation criteria*.
- The leaf node (final level) represent the *result of the classification*. The label (classification decision) for data in each leaf node is made based on the content of the node's data



Information Entropy




Figure 3-3. Entropy of a two-class set as a function of $p(+)$.

- In thermodynamics, entropy represents a level of disorder within a system.
- Similarly, information entropy measure the average level of uncertainty for a random variable. It measure the purity (homogeneity) level of a group of data.
- Information entropy is calculated by the formula:

$$entropy = -p_1 \log(p_1) - p_2 \log(p_2) - \dots$$

$$= \sum_{i=1}^n -p_i \log(p_i)$$
- For 2-classes data, the highest entropy occurs when the data contains samples of both classes in an equal amount as the expected value of a class is at the most uncertain state here.

Entropy

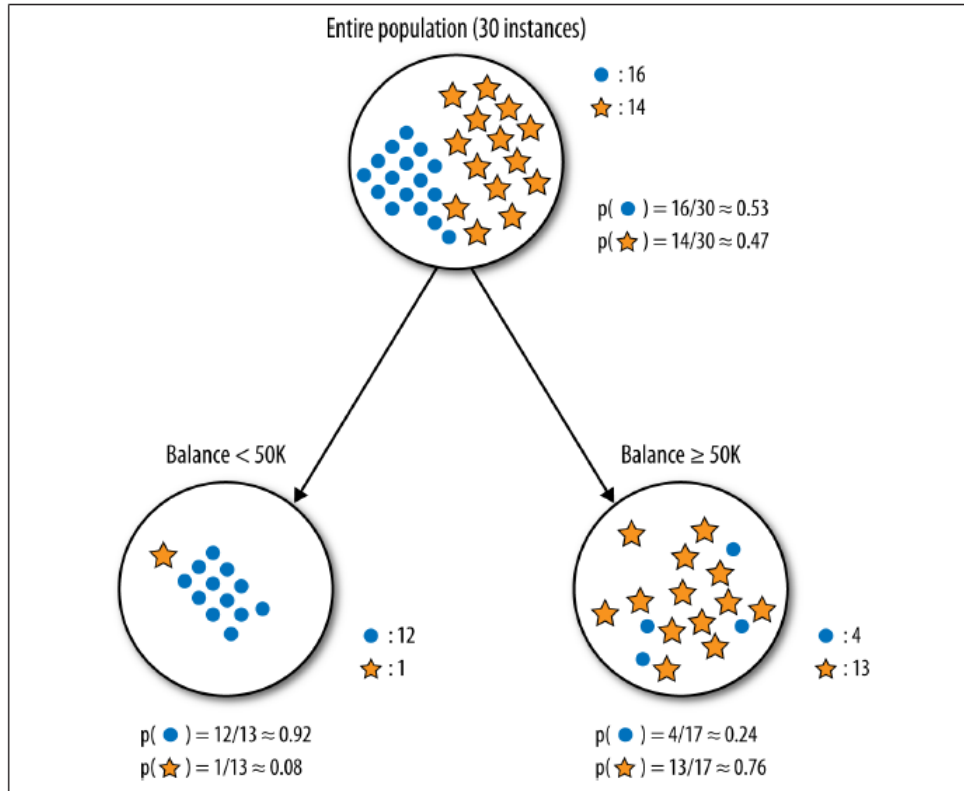


Figure 3-4. Splitting the “write-off” sample into two segments, based on splitting the Balance attribute (account balance) at 50K.

- The entropy of the *parent* (entire population) is

$$\begin{aligned} \text{entropy}(\text{parent}) &= -p(\bullet) \times \log_2 p(\bullet) - p(\star) \times \log_2 p(\star) \\ &\approx -0.53 \times -0.9 - 0.47 \times -1.1 \\ &\approx 0.99 \text{ (very impure)} \end{aligned}$$

- The entropy of the *left* child is:

$$\begin{aligned} \text{entropy}(\text{Balance} < 50K) &= -p(\bullet) \times \log_2 p(\bullet) - p(\star) \times \log_2 p(\star) \\ &\approx -0.92 \times (-0.12) - 0.08 \times (-3.7) \\ &\approx 0.39 \end{aligned}$$

- The entropy of the *right* child is:

$$\begin{aligned} \text{entropy}(\text{Balance} \geq 50K) &= -p(\bullet) \times \log_2 p(\bullet) - p(\star) \times \log_2 p(\star) \\ &\approx -0.24 \times (-2.1) - 0.76 \times (-0.39) \\ &\approx 0.79 \end{aligned}$$

Selecting the Best Split: Information Gain

- The most common splitting criterion is called an *information gain (IG)*
- We would like the split to be *informative*, meaning that it provide us with more information about our classification target. Therefore, we would like for the overall level of entropy to **decrease** after the segmentation
- We call the **reduction in entropy level** (the level of disorder) an *information gain (IG)*
 - i.e. Information Gain = total entropy before split - total entropy after split
- In other words ,

$$\begin{aligned} IG &= \text{entropy}(\text{parent}) - p(c_1) \times \text{entropy}(c_1) - p(c_2) \times \text{entropy}(c_2) - \dots \\ &= \text{entropy}(\text{parent}) - \sum_i p(c_i) \times \text{entropy}(c_i) \end{aligned}$$

- We pick the segmentation criteria to be the one that has the **highest information gain**.

Entropy




Figure 3-4. Splitting the “write-off” sample into two segments, based on splitting the Balance attribute (account balance) at 50K.


- The entropy of the *parent* is 0.99
- The entropy of the *left* child is 0.39
- The entropy of the *right* child is 0.79
- The Information Gain (IG) is:

$$IG = \text{entropy}(\text{parent}) - [p(\text{Balance} < 50K) \times \text{entropy}(\text{Balance} < 50K) + p(\text{Balance} \geq 50K) \times \text{entropy}(\text{Balance} \geq 50K)]$$

$$\approx 0.99 - (0.43 \times 0.39 + 0.57 \times 0.79)$$

$$\approx 0.37$$

Entropy



- $entropy(parent) \approx 0.99$
- $entropy(Residence = OWN) \approx 0.54$
- $entropy(Residence = RENT) \approx 0.97$
- $entropy(Residence = OTHER) \approx 0.98$
- $Information\ Gain\ (IG) \approx 0.13$

Figure 3-5. A classification tree split on the three-valued Residence attribute.

Decision Tree Example: Data and Attributes




Figure 3-2. A set of people to be classified. The label over each head represents the value of the target variable (write-off or not). Colors and shapes represent different predictor attributes.

What are attributes that describes this data (people)?

Attributes:

- head-shape: square, circular
- body-shape: rectangular, oval
- body-color: gray, white

Target variable:

- write-off: Yes, No

What is the most informative attribute that can be used to segment this data (people)?




Figure 3-11. First partitioning: splitting on body shape (rectangular versus oval).

Decision Tree Example: Further Segmentation




Figure 3-12. Second partitioning: the oval body people sub-grouped by head type.




Figure 3-13. Third partitioning: the rectangular body people subgrouped by body color.

Decision Tree Example: Overall Decision Tree


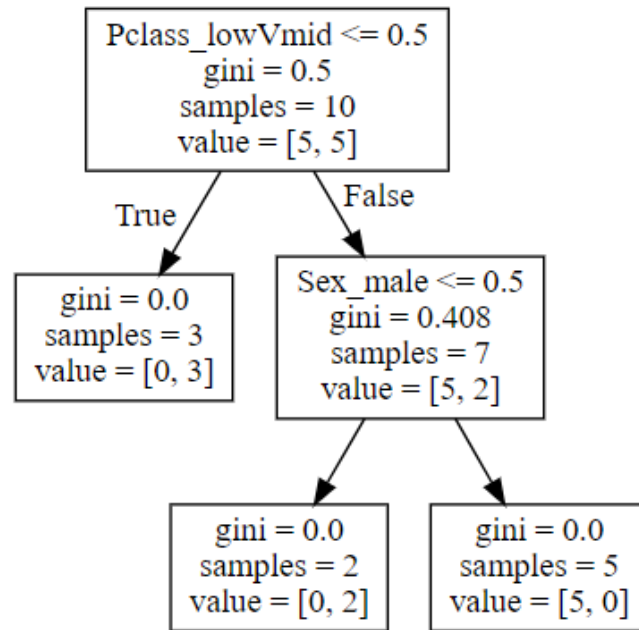


Figure 3-14. The classification tree resulting from the splits done in Figure 3-11 to Figure 3-13.

Note on Decision Tree



$$\text{Gini Index} = 1 - \sum_{i=1}^c (p_i)^2$$


- Decision Tree classification algorithm presented is a *greedy* algorithm.
 - This means it always choose the feature that best segmented the data at any given point.
 - This algorithm can result in sub-optimal solution.
- Another popular criterion used for deciding how to split data at each node of a decision tree is the *Gini impurity*
- Gini impurity is a measurement of the probability that an *incorrectly labeled classification will be given to a randomly chosen element* from a set if the said element is randomly labeled by the distribution of labels in that set. *We want a variable split to have a low Gini index.*

Decision Tree Example


- Colab!

Ensemble Learning

- Ensemble learning is a process which utilizes multiple machine learning models to learn a particular task, then combines their results to make the final prediction.
- For example, an ensemble classifier may take result of its multiple classifiers then deciding the final prediction via voting or some other means of aggregations
- Random Forest and XGBoost models are examples of popular ensembles models



Bagging

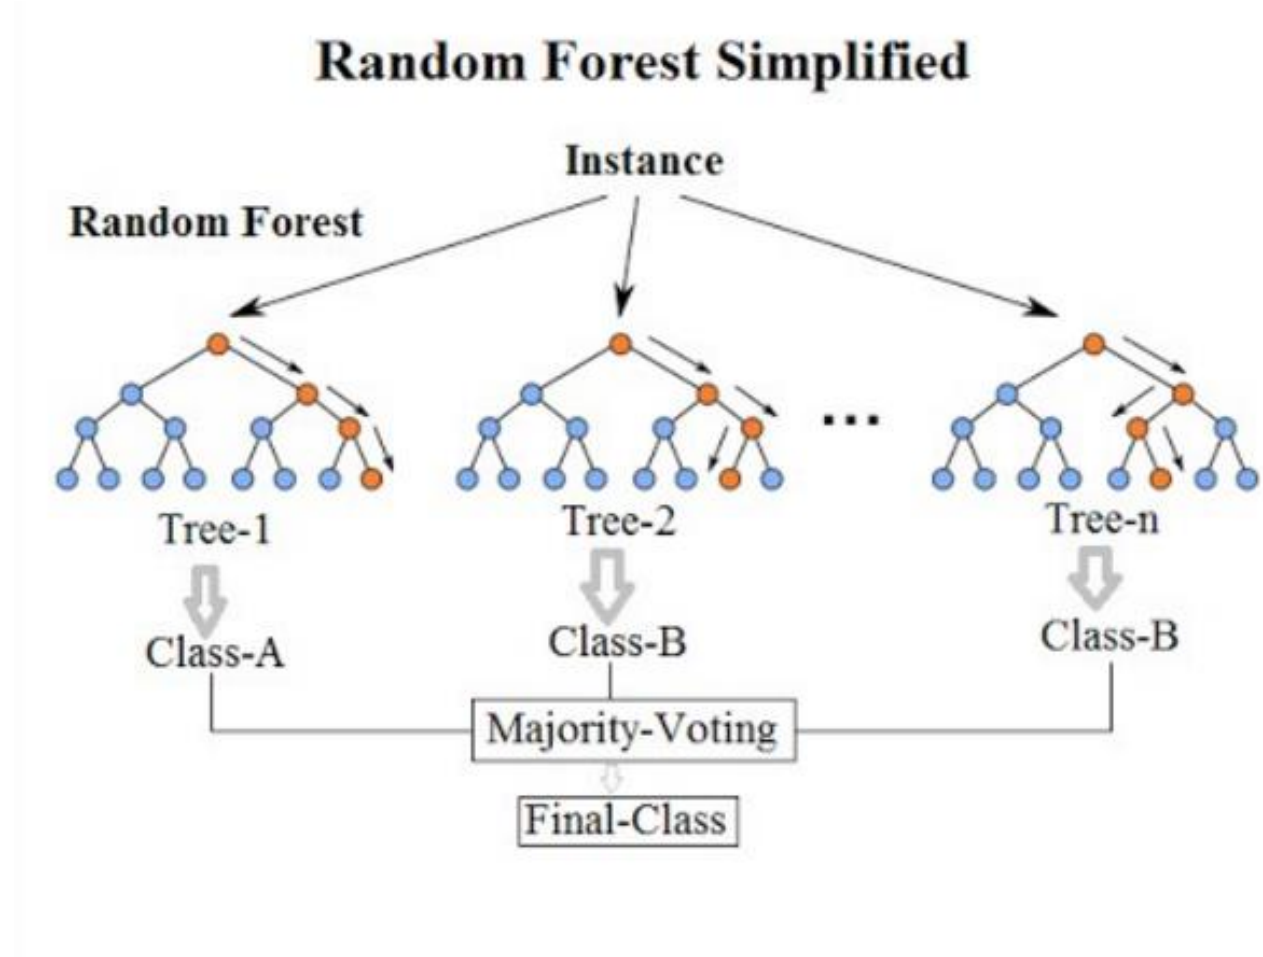


- From the original data, repeatedly perform bootstrapping (taking **random sampling with replacement**) and create multiple sets of bootstrap samples
- Train multiple distinct classifiers on each set of bootstrap samples
- Uses all classifiers to predict a new input by voting / averaging
- Helps making the model more generalized

Random forest

- An ensemble model consists of multiple decision trees, each of which are trained by samples from **bootstrap sampling** with **only a random subset features available**.
 - With large number of trees, the model is less likely to overfit.
 - With only a random subset of features available for each tree, it is less likely for trees' results to be highly correlated.
- Each of decision trees are built and train in a greedy fashion using conditional entropy
- The final prediction is decided by average predictions of all trees.

Random forest



Boosting

- Aims to improve the performance of the model by allowing it to fit the training data better
- Unlike bagging, it **creates a sequence of classifiers** and gives higher influence to a more accurate classifiers
- For each iteration, it makes **examples currently misclassified more important** (or less, in some cases) by giving larger weights in in the construction of the next classifiers.
- The final prediction is decided by combining the results of classifiers by **weighted vote** (weight given by classifier accuracy)



Extreme Gradient Boosting (XGBoost)

- A scalable and portable and accurate implementation of gradient boosting machines
- One of the most popular models. Perform well across variety of applications.

Model Feature

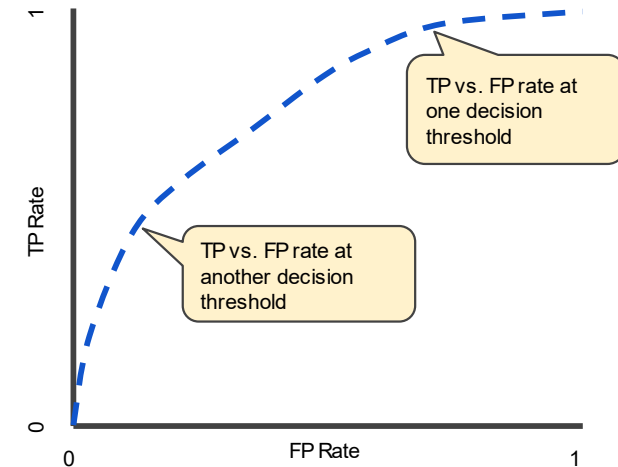
- **Gradient Boosting** algorithm also called gradient boosting machine including the learning rate.
- **Stochastic Gradient Boosting** with sub-sampling at the row, column and column per split levels.
- **Regularized Gradient Boosting** with both L1 and L2 regularization.

Model Selections: ROC curve and AUC

- The ROC (Receiver Operating Characteristic) curve is a plot that shows performance of a model at different classification threshold
- ROC Curves plots Recall vs. Sensitivity, essentially comparing the “hit rate” and the “false alarm rate”

$$\text{True Positive Rate (Recall)} = \frac{TP}{TP + FN}$$


$$\text{False Positive Rate (Sensitivity)} = \frac{FP}{FP + TN}$$



True Label

Predicted Label	True Label	
	Positive	Negative
Positive	True Positive	False Positive
Negative	False Negative	True Negative

Model Selections: ROC curve and AUC



- AUC (Area Under the ROC Curve) provides a metric for considering the models performance across all classification threshold.
- *Side note: so far, we have been using a classification threshold of 0.5. However, it is also possible (and sometimes beneficial) for a classifier to vary.*

Random Forest & XGBoost (& Model Selection) Example

- Colab!



Thank You